



VIT®

Vellore Institute of Technology

(Deemed to be University under section 3 of UGC Act, 1956)

Legal Statute Recommendation for Indian Judgments

BCSE409L - Natural Language Processing

Faculty: Rajeshkannan R

Case Study - III

22BCE3347 – Armaan Indani

22BCE3629 – Nipun Misra

22BCE3469 – Aditya Tripathi

Github Link:

<https://github.com/Armaan-Indani/NLP-Legal-Statute-Identification-ILB.git>

1. ABSTRACT

India's legal system is vast and complex, with millions of cases processed yearly across diverse jurisdictions. A key challenge is the accurate and timely identification of relevant legal statutes (e.g., IPC, CrPC, special Acts) based on case narratives. This task is manual, time-consuming, and even more difficult in regional courts where judgments may be in local languages.

To address this, we propose a Multilingual Legal Statute Recommendation System for Indian Judgments that understands text and audio inputs in English, Hindi, Marathi, Tamil, and Kannada. The system's goal is to automatically recommend the most relevant statutory provisions for a given case description.

It will use state-of-the-art NLP techniques, including transformer-based legal models and speech recognition. At its core is InLegalBERT, a BERT variant fine-tuned on Indian legal texts, combined with advanced statute identification modules trained on datasets like LeSICiN, which link judgments to statutes.

To support multilingual input, models like IndicBERT, MuRIL, and mBERT will be used. A language detection and translation pipeline will ensure consistent semantic interpretation. For audio inputs, ASR models such as Whisper or IndicTTS (trained on Indian legal speech) will enable real-time, spoken-case understanding.

The system will output a ranked list of statutes with confidence scores and explanations, showing the relevance of each suggestion. Additionally, graph-based models may be incorporated to exploit citation networks, improving contextual accuracy based on prior case references.

This tool aims to benefit lower courts and rural legal settings where legal database access and English proficiency are limited. Supporting regional languages and speech input, it promotes accessible legal intelligence, enhancing research efficiency and decision-making.

In summary, this project unites NLP, multilingual AI, legal domain expertise, and speech processing to build an inclusive, intelligent, and scalable Legal Statute Recommendation System

for the Indian judiciary, promoting faster legal aid, linguistic inclusivity, and statutory transparency.

Keywords:

Legal Statute Recommendation, Indian Judiciary, Case Law Analysis, Statutory Interpretation, Legal Research Automation, Multilingual NLP

2. INTRODUCTION

India's legal system is a vast and complex network, rooted in centuries of jurisprudence and evolving through layers of statutes, case law, and constitutional interpretation. It ranks among the largest legal systems globally, with over 25 million cases pending across various levels of the judiciary, including the Supreme Court, High Courts, and a wide network of subordinate courts. Each case is governed by numerous statutes, such as the Indian Penal Code (IPC), Code of Criminal Procedure (CrPC), and Code of Civil Procedure (CPC), along with various special and regional laws. Navigating this system requires not only legal expertise but also the ability to associate specific statutory provisions with the facts, arguments, and context of individual cases.

A core challenge for legal practitioners, judges, researchers, and law students is identifying relevant statutes from case descriptions. While seemingly straightforward, the task is complicated by legal documents written in dense, technical language and by case facts that are often unstructured and vary significantly in presentation. In India, these challenges are amplified by linguistic diversity, uneven access to digital legal databases, and the sheer volume of documentation processed daily [6][7][11]. Manual analysis of lengthy narratives to map them to relevant statutory sections is labor-intensive and prone to omissions or errors, particularly in lower courts or rural settings where digital tools and legal expertise are limited [11][12]. Moreover, many legal proceedings in district and taluka courts are conducted in regional languages, and judgments are often dictated orally or recorded with localized terminology, making reliance on English-only tools inadequate [6][13][15].

To address these gaps, this project proposes a *Multilingual Legal Statute Recommendation System for Indian Judgments*. The system will accept legal case descriptions in both text and speech form, in English and four regional Indian languages, and recommend a ranked list of relevant statutes. Designed as an assistive legal research tool, it aims to help practitioners, students, and citizens quickly identify statutory provisions most relevant to a given case.

At its core, the system leverages advanced Natural Language Processing (NLP) models, specifically transformer-based architectures capable of capturing semantic, syntactic, and contextual nuances [3][4][5]. We employ *InLegalBERT*, a BERT-based model pre-trained and fine-tuned on Indian legal texts such as judgments, statutes, and commentaries [4]. Unlike general-purpose BERT models trained on generic web data, InLegalBERT is tailored to the intricate phrasing and terminology of the Indian legal domain [4][3].

The system also utilizes labeled datasets such as *LeSICiN* (Legal Statute Identification using Citation Network), which maps Indian court judgments to statutes they cite or are influenced by [1]. This enables models that recommend statutes not only through textual similarity but also via legal reasoning and citation relationships are crucial when cases with different surface texts are governed by similar statutes based on precedent or judicial interpretation [1][2][10].

To support multilingual capability, models like *IndicBERT*, *MuRIL*, and *mBERT* are integrated, enabling the system to process inputs in languages such as Hindi, Tamil, Bengali, Marathi, or Kannada [6][15]. A language detection and translation pipeline ensures all inputs are normalized before processing, maintaining semantic consistency regardless of the source language [6][13][15].

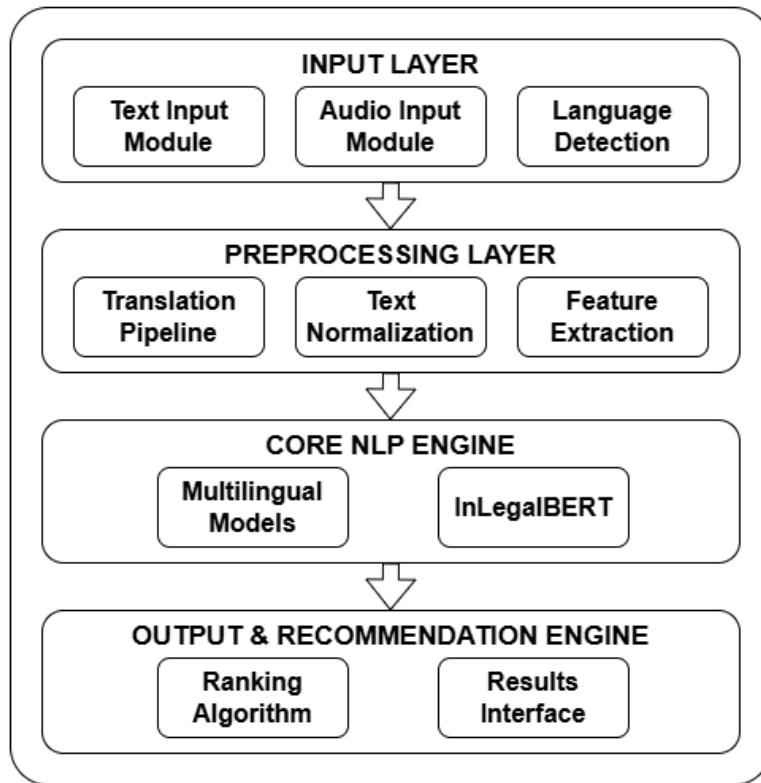
For speech-based inputs, the system incorporates *Automatic Speech Recognition (ASR)* using models like Whisper (by OpenAI) and IndicTTS, trained for Indian accents, code-switching, and legal domain terminology [15]. This is particularly beneficial for users with limited typing skills, disabilities, or in low-literacy regions [6][11].

The output includes a ranked statute list with confidence scores and natural language explanations linking the input to each recommendation [10][15].

The proposed system has diverse applications as a decision-support tool for judges, a research assistant for lawyers and students, a public-facing legal aid service, and an educational resource. Its multilingual and speech-enabled design makes it especially valuable in lower courts, rural jurisdictions, and semi-urban areas where access to legal resources is limited [6][11][13]. By reducing dependency on English-centric systems, it promotes inclusive access to justice [6][7][15].

In conclusion, this *Multilingual Legal Statute Recommendation System* combines advanced NLP, domain-specific modeling, multilingual understanding, and speech recognition to transform statute identification in India. It addresses the technical challenge of automating recommendations while supporting broader goals of legal empowerment, equitable access to justice, and digital transformation of the judicial ecosystem [4][6][15]. This represents a significant step toward intelligent, inclusive, and context-aware legal tools for a linguistically diverse and legally complex nation.

2.1 Proposed Architecture Diagram



3. LITERATURE REVIEW

With the increasing digitalization of court records and the availability of structured legal texts, the intersection of Natural Language Processing (NLP) and legal research has emerged as a powerful domain of applied artificial intelligence. One core task that has received attention in recent years is Legal Statute Identification (LSI), automatically determining the most relevant statutory provisions (e.g., sections of the Indian Penal Code) for a given legal case description. Traditional LSI systems have relied heavily on text classification, but recent advances in deep learning, graph-based networks, and domain-adapted language models have dramatically improved accuracy and scalability. This survey explores five prominent lines of research relevant to building a Multilingual Legal Statute Recommendation System for Indian Judgments, focusing on citation-aware models, graph-augmented methods, and domain-specific language model pretraining.

In their work on Legal Statute Identification (LSI) in Indian law, Paul et al. propose a novel, graph-based model named LeSICIN [1]. They argue that existing LSI methods rely solely on the textual content of legal documents and statutes, ignoring the valuable information available in the legal citation network. To address this, LeSICIN models the documents and statutes as a heterogeneous graph, learning both textual and structural features to predict links between new case documents and relevant statutes. To facilitate this research, the authors curated a large-scale dataset, the Indian Legal Statute Identification (ILSI) dataset, which consists of facts from Indian court cases and a set of 100 statutes from the Indian Penal Code. Experiments showed that LeSICIN significantly outperforms several state-of-the-art baselines, and the ILSI dataset more accurately reflects the multi-label nature of the LSI problem compared to other existing datasets.

A study by Bhattacharya and others proposed a new approach for calculating the similarity between legal case documents by developing a heterogeneous network called Hier-SPCNet [2]. This network augments the traditional Precedent Citation Network (PCNet) by incorporating the hierarchical structure of legal statutes and their citation links with case documents and other statutes. The authors applied the graph embedding algorithm, Metapath2vec, to this network to estimate document similarity. They validated their method on a dataset of 100 pairs of Indian

Supreme Court case documents, whose similarities were annotated by legal experts. The results showed a significant improvement in document similarity estimation compared to methods that only use the PCNet. The study also demonstrated that their network-based method provides complementary insights to a text-based method (Doc2Vec), and that combining the two approaches can be beneficial for estimating legal document similarity.

Chalkidis and colleagues [3] explore the adaptation of BERT, a pre-trained language model, to the legal domain. They investigate three main strategies: using the original BERT model as is, further pre-training it on domain-specific corpora (Legal-BERT-FP), and training a new BERT model from scratch on legal texts (Legal-BERT-SC). The study found that both further pre-training and training from scratch on legal corpora led to better performance on domain-specific tasks compared to using the original BERT model. They also discovered that a wider range of hyperparameter tuning during the fine-tuning phase can lead to substantial performance improvements. The authors introduced and released a family of legal BERT models, including a smaller, more efficient version called Legal-BERT-SMALL, which was shown to be competitive with larger models despite its smaller size.

Building on this, InLegalBERT [4] was trained on 5.4 million Indian legal documents and demonstrated superior performance across tasks like statute identification, judgment segmentation, and appeal outcome prediction. The model outperformed LegalBERT and CaseLawBERT on both Indian and EU/UK datasets, affirming the benefit of country-specific pretraining and custom vocabularies for legal tasks.

A paper by Dhanani et al. addresses the need for an effective and scalable legal judgment recommendation system (LDRS) for legal professionals [5]. The authors propose a pre-learned word embedding-based LDRS (P-LDRS) that uses domain-specific pre-learned word embeddings to initialize the Doc2Vec model, thereby enhancing the semantic representation of legal judgments. To tackle the scalability issues associated with large and growing legal document corpora, they also introduced a distributed version of their P-LDRS that leverages frameworks like MapReduce and Spark. Their empirical analysis on a dataset of over 48,000 Indian Supreme Court judgments demonstrates that the non-distributed P-LDRS performs significantly better than a traditional Doc2Vec-based LDRS, achieving an Accuracy of 0.88, an

F1-score of 0.82, and an MCC Score of 0.73. Furthermore, the distributed P-LDRS shows improved time efficiency with a stable performance, confirming its effectiveness and scalability for handling large volumes of legal text.

The application of Natural Language Processing (NLP) within the legal sector has experienced substantial growth, motivated by the necessity to manage intricate legal documents and optimize judicial workflows. Ariaai et al. [6] provide an extensive survey detailing the progression of legal NLP, systematically categorizing key tasks such as legal document summarization, named entity recognition, legal question answering, judgment prediction, and argument mining. Their analysis underscores challenges unique to legal language, including the considerable length of documents, scarcity of structured datasets, and concerns related to fairness and explainability in legal contexts.

Addressing these challenges in the Indian legal environment, Joshi et al. [7] introduced IL-TUR, a multilingual benchmark designed specifically for Indian legal text comprehension and reasoning. IL-TUR encompasses eight distinct tasks ranging from rhetorical role labeling to legal summarization across English and nine Indian languages. This benchmark not only provides baseline performance metrics but also highlights the limitations of existing large language models (LLMs) when applied to authentic Indian legal texts, which are frequently lengthy, unstructured, and domain-specific.

In an endeavor to improve legal information retrieval within the Indian legal framework, Pandian and Joshi [8] introduced a context-aware recommendation system for statutes and case law, leveraging transformer-based masked language models. Their research underscores the superiority of semantic representations over traditional lexical matching approaches, particularly in the retrieval of pertinent statutes and judicial precedents. By fine-tuning language models originally developed for general domains on Indian legal corpora, their methodology demonstrates a marked enhancement in the accuracy of statute recommendations relative to conventional information retrieval techniques. Importantly, the study identifies a significant challenge: while statutes can be retrieved with greater consistency through semantic similarity measures, recommending prior case law proves more complex due to its inherent intricacy and heterogeneity. This finding highlights the necessity for further domain-specific adaptation and

the development of specialized modeling strategies tailored to the Indian legal natural language processing context.

Concurrently, statute recommendation systems have emerged as valuable tools for legal practitioners. Feng et al. [9] proposed a portable, two-stage neural network-based statute recommender incorporating attention mechanisms and a ranking module. This framework advances multi-label statute prediction by leveraging deep learning techniques to capture contextual and semantic relationships among legal provisions, thereby improving prediction accuracy.

Complementing this approach, Li et al. [10] developed a more interpretable methodology employing hand-crafted relational features to model the relationship between cases and statutes. Their hybrid system integrates collaborative filtering for initial candidate retrieval with feature-based learning-to-rank models for final statute selection. The incorporation of explainable features enhances interpretability and practical applicability, particularly in civil law domains such as divorce cases in China.

Deshmukh and Kamble [11] developed IndianBailJudgments-1200, a structured dataset consisting of 1,200 Indian bail orders annotated via prompt-engineered GPT-4o. It provides fields for fairness analysis, judgment prediction, and social research. While the dataset is pioneering, its focus on High Court judgments and English-only documents limits its generalization.

Kalamkar et al. [12] focused on Named Entity Recognition (NER) within legal texts, presenting a dataset with 46,545 annotated entities from Indian court judgments. Baseline models like RoBERTa, supported by post-processing for coreference and long-span handling, achieved high accuracy. However, sentence-level annotations and English-only judgments pose constraints.

Joshi et al. [13] introduced INDICourtSumm, a large-scale summarization dataset with over 31,000 Indian legal cases and summaries derived from IndianKanoon. Fine-tuned PEGASUS and BART models showed strong summarization performance (ROUGE-1: ~57.4), but the lack of gold-standard manual summaries introduces evaluation noise.

Malik et al. [14] contributed the ILDC for the CJPE dataset, designed for judgment prediction and explanation tasks. Using models like BiLSTM and Legal-BERT, the study achieved over 81% F1-score in binary outcome classification and provided rationale-based interpretability. Despite its innovation, the explanation depth and case complexity remain limited.

Bhattacharya and Ghosh [15] compared extractive and abstractive summarization approaches on Indian court case documents. While extractive methods like TextRank and LexRank performed reasonably well, abstractive models such as T5 and BART generated more concise and coherent summaries. The work highlighted challenges like domain-specific terminology handling and document length limitations that hinder model accuracy.

Collectively, these studies contribute crucial resources and insights into the growing field of Legal NLP in India. The ongoing challenges emphasize the need for multilingual, multi-tier court data, advanced reasoning models, and better factual consistency in generated outputs.

3.1 Summary Table

S. No	Title and Authors (Year)	Methodology / Approach & Dataset	Key Findings / Contributions	Limitations / Gaps
1	LeSICiN: A Heterogeneous Graph-based Approach for Automatic Legal Statute Identification from Indian Legal Documents - Paul et al. (2022)	Uses a heterogeneous graph to model relationships between case facts and statutes. Treats Legal Statute Identification (LSI) as an inductive link prediction problem. Trained on a new Indian dataset mapping court judgments to IPC sections.	Achieved 19.2% improvement over baselines. Handles unseen test cases effectively. Introduced a new benchmark dataset for Indian LSI.	Limited to IPC statutes only. Focused solely on textual judgments; lacks multimodal support (e.g., speech).

	Hier-SPCNet: A Legal Statute Hierarchy-based Heterogeneous Network for Computing Legal Case Document Similarity - Bhattacharya et al. (2020)	Constructs a heterogeneous network combining case citations and statute hierarchy (Act → Chapter → Section). Applies MetaPath2Vec for node embeddings. Tested on the Indian court case corpus.	Demonstrated improved document similarity estimation by leveraging statute structure. Enhanced legal case retrieval and clustering accuracy.	Focuses on case similarity, not statute identification. Does not include multilingual or audio inputs.
2	LEGAL-BERT: The Muppets Straight Out of Law School - Chalkidis et al. (2020)	Compares three strategies: using generic BERT, domain-adaptive pretraining, and training BERT from scratch on legal text. Evaluated on European legal corpora (EU/UK) across multiple legal NLP tasks.	Domain-pretrained models significantly outperform off-the-shelf BERT. LEGALBERT models achieve SOTA on classification and QA tasks in legal settings.	Not trained on Indian legal texts. Vocabularies and context may not generalize well to Indian law.
3	Pre-trained Language Models for the Legal Domain: A Case Study on Indian Law - Paul et al. (2023)	Developed and compared three Indian legal PLMs: InLegalBERT, InCaseLawBERT, and CustomInLawBERT. Pretrained on 5.4 million Indian legal documents from Indian Kanoon. Evaluated on 5 tasks: statute identification, segmentation, outcome prediction, etc.	InLegalBERT outperforms other PLMs (including LegalBERT) on 4/5 tasks. Shows the value of country-specific corpora and vocabularies. Introduces new Indian legal benchmarks.	Does not yet support regional Indian languages. Focuses only on text, lacks audio or multimodal input support.

5	<p>Effective and Scalable Legal Judgment Recommendation Using Pre-learned Word Embedding - Dhanani et al. (2022)</p>	<p>Uses Doc2Vec and domain-specific embeddings (e.g., Law2Vec, LegalW2V). Designed a scalable variant using Apache Spark and MapReduce for large corpora. Trained on Indian court judgments from public legal repositories.</p>	<p>Improves the accuracy and F1score of judgment recommendation. Demonstrates scalable training for large datasets. Addresses sparse citation limitations using semantic embedding.</p>	<p>Embedding-based models lack interpretability. Does not utilize graph structure or deep contextual models (e.g., BERT).</p>
6	<p>IL-TUR: Benchmark for Indian Legal Text Understanding and Reasoning - Abhinav Joshi et al., 2023</p>	<p>Developed IL-TUR benchmark covering 8 legal NLP tasks; datasets in English and 9 Indian languages; baseline models using LLMs and traditional methods</p>	<p>Introduced the first comprehensive Indian legal NLP benchmark; public leaderboard for legal text understanding and reasoning; highlighted LLM performance gaps</p>	<p>Dataset limited to the Indian legal system; large document length is still challenging for LLMs; scope for adding more tasks</p>
7	<p>Natural Language Processing for the Legal Domain: A Survey of Tasks, Datasets, Models and Challenges - Farid Ariai et al., 2025</p>	<p>Systematic review of 131 papers across legal NLP tasks: NER, Summarisation, Question Answering, Argument Mining, Judgement Prediction, etc.</p>	<p>Provided a broad overview of the legal NLP landscape, tasks, datasets, and models; identified 16 research challenges, including bias, explainability, and data scarcity</p>	<p>Survey only; no new dataset or model; primarily focuses on English language works; some sub-tasks are underexplored</p>

8	<p>Autosuggestion of Relevant Cases and Statutes</p> <p>Pandian & Joshi (2022)</p>	<p>Utilized BERT-based masked language models to recommend relevant legal statutes and case citations.</p> <p>Dataset: A curated collection of Indian legal queries and associated citations/statutes (details not fully public).</p>	<p>Semantic retrieval significantly outperformed lexical baselines for statute recommendation.</p> <p>Statutes were easier to retrieve accurately than prior case laws.</p>	<p>The approach shows weaker performance in recommending past cases.</p> <p>Dataset availability and generalizability to real court queries remain limited.</p>
9	<p>Recommendng Statutes: A Portable Method Based on Neural Representations - Yi Feng et al., 2020</p>	<p>Two-step neural network system: RNN + Max Pooling + Attention + Statute2Vec embedding; trained on Chinese legal datasets</p>	<p>Improved statute recommendation accuracy over traditional baselines; portable method for multi-label statute prediction</p>	<p>Dataset limited to Chinese court data; method ignores potential legal domain knowledge graphs</p>
10	<p>Statute Recommendation: Re-ranking statutes by modeling case-statute relation with interpretable hand-crafted features - Chuanyi Li et al., 2022</p>	<p>Used collaborative filtering to retrieve candidate statutes; applied hand-crafted relational features with pairwise ranking models; tested on Chinese legal data</p>	<p>Achieved ~5% recall improvement over collaborative filtering; an interpretable feature set helps understand the case-statute relation</p>	<p>Focused on divorce cases in Chinese courts; not generalized to broader legal domains or jurisdictions</p>

11	IndianBailJudgments-1200: A Multi-Attribute Legal NLP Dataset for Bail Order Understanding in India Deshmukh & Kamble, 2025	Dataset of 1200 Indian bail judgments annotated using prompt-engineered GPT-4o; validated subset manually; 20+ structured fields designed via schema for legal NLP	First dataset focused on Indian bail jurisprudence; enables tasks like fairness analysis, legal summarization, judgment prediction; supports legal, social, and AI research	Limited to High Court decisions; annotations are partly subjective; relies heavily on LLMs; English-only cases; lacks lower court coverage
12	Named Entity Recognition in Indian court judgments Prathamesh Kalamkar et al., 2022	Created a legal NER corpus of 46,545 annotated entities from Indian court judgments (14 legal entity types); trained baseline NER models (Roberta + Transition Parser & Fine-Tuned Transformers); dataset includes Supreme & High Court judgments	Developed legal NER baseline models with high F1 (91.1%); introduced post-processing techniques for coreference resolution and document-level corrections; released datasets and models publicly	Sentence-level annotation limits context; issues in long entity spans (e.g., precedents); English-only judgments; biases possible due to source sampling and heuristic case classification
13	INDICourtSumm: A Dataset for Judgment Summarization of Indian Court Cases Abhishek Joshi et al., 2023	Created INDICourtSumm with ~31,000 Indian court judgments and summaries from IndianKanoon; explored summarization using PEGASUS and BART models with legal-domain adaptations	First large-scale summarization dataset for the Indian legal domain; PEGASUS fine-tuned on INDICourtSumm outperformed baselines (ROUGE-1 up to 57.4)	The dataset has noisy summaries from IndianKanoon; it lacks gold-standard manual summaries; is focused only on extractive

				summaries (not abstractive generation or reasoning)
14	ILDC for CJPE: Indian Legal Documents Corpus for Court Judgment Prediction and Explanation Malik et al., 2021	<p>Introduced ILDC corpus with 35K Indian court cases (Supreme and High Courts); annotated for judgment outcome; built classifiers (SVM, BiLSTM, BERT) for judgment prediction and explanation</p>	<p>Proposed benchmark dataset for CJPE tasks in Indian law; fine-tuned Legal-BERT achieved 81% F1 for binary outcome prediction; introduced rationale-based explanation via attention weights</p>	<p>Focused on outcome classification only; explanation limited to token-level attention (may not reflect true rationale); lacks detailed case reasoning or multi-label outcomes</p>
15	Legal Case Document Summarization: Extractive and Abstractive Methods and Their Evaluation, Bhattacharya et al., 2022	<p>Utilized transformer models like BERT, GPT-2, and LED for both extractive and abstractive summarization. Benchmarked using the Indian Legal Documents Corpus (ILDC) and INTACT datasets.</p>	<p>Introduced a comprehensive evaluation of summarization techniques on legal data. Found that the LED (Longformer) performed well for long documents. Released benchmark datasets and evaluation results.</p>	<p>Abstractive models still struggle with factual consistency; performance varies significantly with document length and type. Limited multilingual evaluation.</p>

4. PROBLEM DESCRIPTION

India's judiciary is one of the largest in the world, with millions of cases pending across multiple levels of courts. It faces a critical challenge in mapping unstructured case descriptions to relevant statutory provisions. The process is currently manual, time-consuming, and highly dependent on legal expertise. This challenge is intensified in:

- **Multilingual settings**, where court proceedings and judgments are delivered in regional languages (Hindi, Marathi, Tamil, Kannada, etc.), which English-only tools cannot process effectively.
- **Speech-based proceedings**, where oral dictations and recordings need transcription and interpretation before statute identification.
- **Resource-limited courts**, especially in rural areas, where access to digital legal databases and English-trained professionals is scarce.

Thus, there is a critical need for an automated, multilingual, and speech-enabled Legal Statute Recommendation System that can:

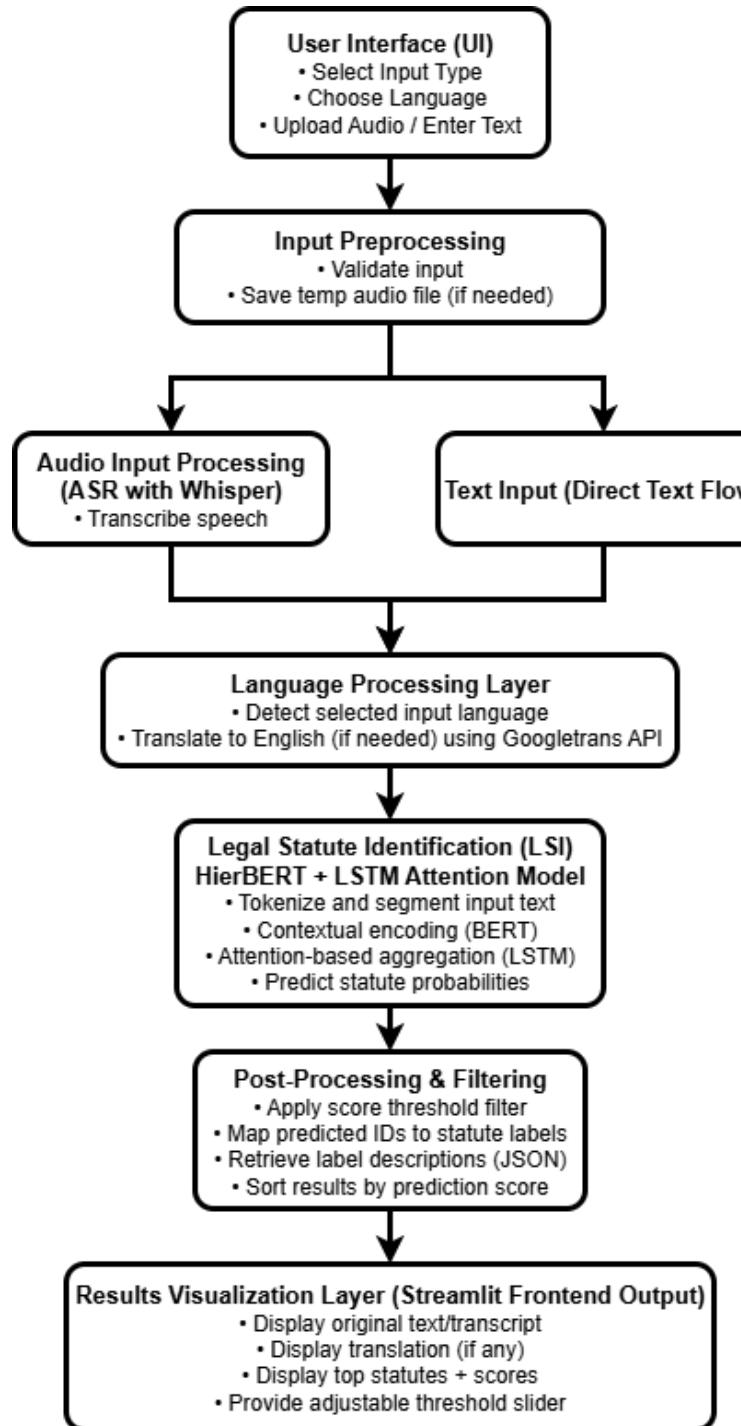
- Understand case descriptions in multiple Indian languages (both text and audio).
- Recommend the most relevant statutory provisions with confidence scores.
- Support judges, lawyers, students, and citizens in quickly identifying applicable laws.

To address this, the proposed system integrates three key components:

1. **Audio-to-Speech Processing** – Converting spoken case narratives into accurate text transcripts using ASR models like Whisper/IndicTTS.
2. **Multilingual Translation & Normalization** – Detecting input language and ensuring semantic consistency across English and Indian languages using models like IndicBERT, MuRIL, and mBERT.
3. **Legal Statute Identification** – Leveraging domain-specific NLP models (InLegalBERT, LeSICiN) to automatically recommend the most relevant statutes with ranked outputs and confidence scores.

This pipeline enables automated, multilingual, and speech-enabled statute recommendation, reducing delays, enhancing accessibility, and supporting practitioners, students, and citizens in navigating India's complex legal system.

4.1 Framework Diagram



4.2 Pseudo Code of Proposed System

```
BEGIN

    IMPORT required libraries
        (Streamlit, Torch, Whisper, Transformers, etc.)

    DEFINE language mappings and translator

    CLASS LstmAttn:
        FUNCTION forward():
            Apply bidirectional LSTM with attention mechanism

    CLASS HierBert:
        FUNCTION forward():
            Encode input segments using BERT
            Aggregate segment representations via LSTM with attention

    CLASS HierBertForTextClassification:
        FUNCTION forward():
            Encode text hierarchically
            Apply classification head with sigmoid activation

    FUNCTION load_lsi_model_and_tokenizer():
        Load label vocabulary, tokenizer, and pretrained model
        Initialize hierarchical BERT and classification head
        Load fine-tuned model weights
        RETURN model, tokenizer, id2label mapping

    FUNCTION load_label_descriptions():
        Load JSON file containing label descriptions
        RETURN dictionary of label descriptions

    FUNCTION load_asr_pipeline():
        Load Whisper ASR model for automatic speech recognition
        RETURN ASR model

    FUNCTION transcribe_audio(audio_file, asr_model, language_code):
        Save audio file temporarily
        Call asr_model.transcribe(language_code)
        Concatenate transcript segments
        RETURN transcribed text

    FUNCTION translate_to_english_googletrans(text):
        Translate text to English using Google Translator API
```

```

    RETURN translated text

FUNCTION get_predictions(text, model, tokenizer, id2label, label_dict,
threshold):
    Split text into sentences
    Tokenize input
    Pass through model to obtain logits
    Apply sigmoid activation and filter scores above threshold
    Sort results by confidence score
    RETURN top 10 predicted statutes with labels

FUNCTION main():
    Initialize Streamlit web interface
    Load all models and required resources
    Display sidebar parameters (threshold, input type, language)
    Obtain input (text or audio)

    IF input is audio:
        Transcribe audio to text
    IF input language ≠ English:
        Translate text to English

    Generate predictions using model
    Display output results as formatted table

CALL main()

END

```

4.3 Explanation of Logic

- The system starts with raw input (text/audio).
- If it's audio, ASR converts it into text.
- Language detection and translation ensure that all inputs are standardized.
- Preprocessing normalizes legal text.
- InLegalBERT extracts statute relevance.
- LeSICiN graph boosts accuracy via citation relationships.
- Final output = ranked statutes with scores.

4.4 Flow Diagram

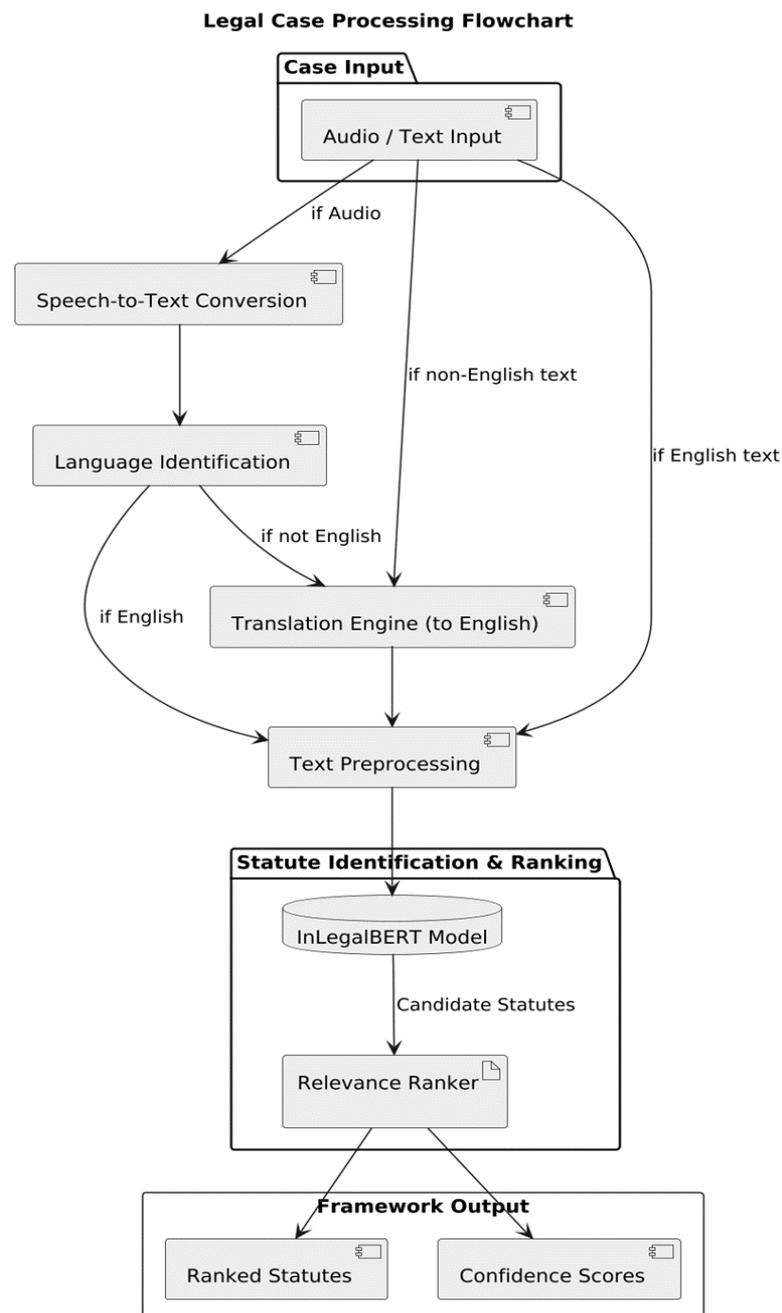


Figure 3. Stepwise pipeline of the Legal Statute Identification framework.

1. Case Input

- Audio / Text Input:
 - The system accepts case descriptions either as spoken audio (e.g., courtroom dictation, lawyer narration) or as typed text (e.g., judgment documents).
 - This ensures flexibility, especially for rural or regional courts where proceedings may not always be recorded in English text.

2. Speech-to-Text Conversion (ASR)

- If the input is audio, it passes through an Automatic Speech Recognition (ASR) system such as Whisper or IndicTTS.
- The ASR converts spoken language (Hindi, Tamil, Marathi, Kannada, etc.) into raw text.
- This step is crucial for handling oral dictations, interviews, and voice-based inputs.

3. Language Identification

- Once text is available (either directly from input or via ASR), the system uses a language detection model to identify the input language.
- This helps the pipeline decide whether translation is needed.
- Example: If the case text is in Hindi, the system knows it must translate it into English before statute identification.

4. Translation Engine (to English)

- If the detected language is not English, the system uses translation models (IndicBERT, MuRIL, mBERT) to convert the text into standardized English.
- If the text is already in English, this step is skipped.
- Ensures all case narratives are semantically consistent before passing them to the statute identification stage.

5. Text Preprocessing

- Raw text often contains noise, repetitions, or informal phrasing.
- Preprocessing includes:
 - Removing stopwords and irrelevant tokens.
 - Normalizing legal terms (e.g., “IPC Sec. 307” → “Section 307 IPC”).
 - Handling spelling and formatting differences.
- Produces a clean, normalized case description ready for statute analysis.

6. Statute Identification & Ranking

This is the core of your system, consisting of two submodules:

1. InLegalBERT Model

- A transformer-based legal NLP model, fine-tuned on millions of Indian legal documents.
- Extracts the semantic meaning of the case description and generates candidate statutes that may apply.

2. Relevance Ranker (LeSICiN Graph / Citation Network)

- Uses legal citation networks to refine the recommendations.
- Ensures statutes aren't just textually similar but also contextually and legally relevant (based on how statutes are cited in past cases).
- Ranks statutes in order of importance.

7. Framework Output

- The final output is presented as:
 - **Ranked Statutes:** Ordered list of the most relevant legal provisions (e.g., IPC Sec 302, CrPC Sec 164).
 - **Confidence Scores:** Probabilities indicating how strongly each statute applies.

- **Explanations:** A natural language justification for why the statute is suggested, improving interpretability.

5. EXPERIMENTS

5.1. Dataset

The project uses three types of datasets, each aligned with a module in the framework:

1. Speech-to-Text Dataset

Source: ULCA ASR Dataset Corpus

Link: <https://github.com/Open-Speech-EkStep/ULCA-asr-dataset-corpus>

Parameters

- audio_id: Unique identifier for each audio file.
- audio_file: Path/link to the audio file (speech recording).
- Language: Language of the audio (Hindi, Marathi, Tamil, Kannada, etc.).
- transcript: Ground truth text transcription of the audio.
- duration: Length of the audio file in seconds.
- speaker_id (if available): Identifier for the speaker (for voice diversity).

Sample Dataset

audio_id	audio_file	language	transcript	duration	speaker_id
A001	hindi_case1.wav	Hindi	आरोपी ने चोरी की घटना की...	12.4	S001

A002	marathi_case2.wav	Marathi	प्रतिवादीने गंभीर इजा केली...	10.1	S002
A003	tamil_case3.wav	Tamil	குற்றவாளி காயப்படுத்தினார் ...	14.8	S003

Explanation

This dataset provides multilingual audio recordings with corresponding human-verified transcripts. It is used to train and evaluate ASR models (Whisper, IndicTTS) for accurate speech-to-text conversion in multiple Indian languages.

2. Translation Dataset

Source: Hugging Face Translation Datasets (based on OPUS, Flores, etc.)

Link: <https://huggingface.co/docs/transformers/en/tasks/translation>

Parameters

- sentence_id: Unique identifier for each example.
- source_language: Language of the input text.
- target_language: Translation target (English).
- source_text: Original text in source language.
- translated_text: Human-translated equivalent in target language.

Sample Dataset

sentence_id	source_language	target_language	source_text	translated_text
T001	Hindi	English	आरोपी ने चोरी की घटना की...	The accused committed theft...
T002	Tamil	English	குற்றவாளி காயப்படுத்தினார் ...	The accused caused injury...
T003	Kannada	English	ಆರೋಪಿ ಹಲ್ಲಿ ನಡೆಸಿದನು...	The accused carried out an assault...

Explanation

This dataset contains parallel sentences in Indian languages and English. It is used to train/fine-tune IndicBERT, MuRIL, and mBERT for multilingual translation, ensuring all case inputs are normalized into English before legal analysis.

3. Legal Statute Identification Dataset (LeSICiN)

Source: LeSICiN: Legal Statute Identification using Citation Networks

Link: <https://zenodo.org/records/6053791>

Parameters

- `case_id`: Unique identifier for a case document.
- `case_text`: Case description or judgment text.

- citations: Statutes or prior cases cited in the judgment.
- statutes_applicable: Ground truth statutes applied (labels).
- predicted_statutes (for evaluation): Model output of recommended statutes.

Sample Dataset

case_id	case_text	citations	statutes_applicable
C001	The accused broke into the house and stole	IPC Sec 378, 411	IPC Sec 378, Sec 411
C002	The husband abandoned and tortured his wife	IPC Sec 498A	IPC Sec 498A
C003	The defendant attacked with a deadly weapon	IPC Sec 307	IPC Sec 307

Explanation

This dataset links Indian court judgments to statutes they cite or are influenced by. It is graph-based, combining textual content with citation networks. This allows the system to recommend statutes not only by semantic similarity but also by legal reasoning and precedent.

Combined Dataset Role in Project

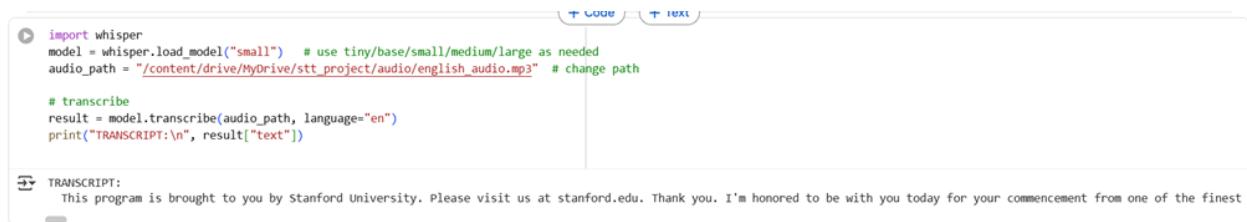
- **ULCA ASR Dataset** → Used for Audio-to-Text Conversion.
- **Translation Dataset** → Used for Multilingual Normalization.
- **LeSICiN Dataset** → Used for Statute Identification & Ranking.

Together, these datasets cover the entire pipeline: speech input → translation → statute recommendation.

6. RESULTS AND DISCUSSION

6.1 Google Colab

Transcription Output



```
import whisper
model = whisper.load_model("small") # use tiny/base/small/medium/large as needed
audio_path = "/content/drive/MyDrive/stt_project/audio/english_audio.mp3" # change path

# transcribe
result = model.transcribe(audio_path, language="en")
print("TRANSCRIPT:\n", result["text"])
```

TRANSCRIPT:
This program is brought to you by Stanford University. Please visit us at stanford.edu. Thank you. I'm honored to be with you today for your commencement from one of the finest

Figure 4: Speech-to-text transcription using the Whisper model in Python. The model successfully transcribes an English audio clip into text.



```
from faster_whisper import WhisperModel
model = WhisperModel("small", device="cuda", compute_type="float16")
audio_path = "/content/drive/MyDrive/stt_project/audio/hindi_audio.mp3"
segments, info = model.transcribe(audio_path, language="hi")
transcript = " ".join([seg.text for seg in segments])
print("Hindi transcript:\n", transcript)
```

Hindi transcript:
याद कर वल्कुपन में एक खेल खेला है तो सन रेस के उपर उस टींस को फोकस करते थे और उस टींस के बीच में से सन रेस निकलती थी पहली भीम की तरा और नीचे माचिस की तीव्री रखके तो जब सन की रेस को फोकस

Figure 5: Transcription of Hindi audio using the Faster Whisper model. The model processes an audio clip in Hindi and outputs the corresponding text transcription.

Translation Output

		Text	English_Translation
0		सुप्रभात, आप कैसे हैं?	Good morning, how are you?
1		शुभ प्रभात, तुम्ही कसे आहात?	Good morning, how are you?
2	का	लेल वणेक्कम, नींनकला एप्पाटि इरुक्किऱ्यकला?	Good morning, how are you?
3		ಶುಭೋದಯ, ನೀವು ಹೇಗೆದ್ದೀರಿ?	Good morning, how are you?

Figure 6: Multilingual text samples and their corresponding English translations demonstrating the language translation module's functionality.

Legal Statute Identification Output

Case 1: Tuliya Devi, married to Gullu about 5 years prior, was subjected to cruelty and harassment for dowry by her husband Gullu, father-in-law Bal Chand, and mother-in-law Ramwati. On Predicted Statutes (Re-ranked): 1. IPC 304B: Dowry death within seven years of marriage. (adjusted_score=1.1607) 2. IPC 494: Marrying again during lifetime of husband or wife (bigamy). (adjusted_score=0.6668)
Case 2: In connection with Raiganj Police Station Case No. 994 of 2013 (under sections 498A/34 IPC), the petitioners applied for anticipatory bail. The Court, after hearing both parties, Predicted Statutes (Re-ranked): 1. IPC 229A: Failure to appear in court after release on bail or bond. (adjusted_score=0.6015) 2. IPC 34: Liability of all persons for a criminal act done with common intention. (adjusted_score=0.4861)

Figure 7: Predicted legal statutes for sample case texts after re-ranking using the proposed Legal Statute Identification (LSI) model.

6.2 GUI Interface

English Text

The screenshot shows the 'Model Parameters' section with a 'Prediction Score Threshold' slider set at 0.50. The main area is titled 'Legal Statute Identification (LSI)' with the subtitle 'Predict relevant sections of The Indian Penal Code from text or audio input.' Below this, the first step is '1. Choose Input Type and Language'. It includes a 'Select Input Type' dropdown with 'Text Input' (selected) and 'Audio Input (File Upload)', and a 'Select Input Language' dropdown set to 'English'. The second step is '2. Provide Input Data', which contains a text area with a case summary about a college student named Neeraj Sharma who altered a vehicle's license plate and sold it to a local mechanic using a forged registration certificate. The summary details his initial denial and subsequent confession after digital payment records and CCTV footage were presented. A note at the bottom of the text area states: 'Flow: English Text → Statute Prediction'. At the bottom right is a red 'Identify Statutes' button.

The screenshot shows the 'Model Parameters' section with a 'Prediction Score Threshold' slider set at 0.50. The main area is titled 'Identify Statutes' with the subtitle 'Step 3: Identifying Statutes...'. Below this, the results are displayed under the heading 'Results'. It shows 'Input Language: English | Input Type: Text Input'. The 'Statute Prediction Output' section displays a table titled 'Found 4 relevant statute sections (Score > 0.50):'. The table has columns for 'Statute Section', 'Prediction Score', and 'Description'. The rows are:

	Statute Section	Prediction Score	Description
0	Section 471 in The Indian Penal Code	0.6473	Using as genuine a forged document or electronic record - Whoever fraudulently or dishonestly uses as genuine a forged document or electronic record -
1	Section 420 in The Indian Penal Code	0.6445	Cheating and dishonestly inducing delivery of property - Whoever cheats and thereby dishonestly induces another person to deliver property -
2	Section 468 in The Indian Penal Code	0.6304	Forgery for purpose of cheating - Whoever commits forgery, intending that the document or electronic record so forged shall be used for cheating -
3	Section 467 in The Indian Penal Code	0.5729	Forgery of valuable security, will - Whoever forges a document which purports to be a valuable security or a

Figure 8. Workflow demonstration of the Legal Statute Identification (LSI) model. (Top) A user provides a case summary in natural language as text input. (Bottom) The system analyzes the text and returns a ranked list of predicted statutes, including sections 471, 420, 468, and 467 of the Indian Penal Code, based on the input summary.

Non-English text

The screenshot shows the LSI system's user interface for handling non-English input. On the left, a sidebar titled "Model Parameters" includes a "Prediction Score Threshold" slider set to 0.50. The main content area has a title "Legal Statute Identification (LSI)" and a subtitle "Predict relevant sections of The Indian Penal Code from text or audio input." Below this, a section titled "1. Choose Input Type and Language" shows "Text Input" selected. To the right, another section titled "2. Provide Input Data" contains a text input field with Hindi text. A large callout box over the text input provides a summary of the case in English:

अमित सिंह, 35 वर्ष का व्यक्ति, अपने पड़ोसी रमेश के साथ जमीन के विवाद को लेकर लंबे समय से झगड़े में था। दोनों के बीच कई बार पंचायत में समझौते की कोशिश हुई, लेकिन विवाद सुलझ नहीं पाया। कुछ दिन पहले दोनों के बीच फिर से बहस हुई, जो हाथापाई में बदल गई। गवाहों के अनुसार, अमित ने मुख्य में आकर धर से लोटे की रोड उठाई और रमेश के सिर पर कई बार चार किया। रमेश भी पर कई बार चार किया। घटनास्थल से बरगद टोंक पर सराश का खुन पाया गया, और अधिकतम का कपड़ा से भी मिलते जुलते खुन के थब्बे मिले। गवाहों के बयान और फारंसिक साक्ष्य ने हत्या का संष्टान प्राप्त किया। यदि अदालत में अमित को आजीवन कारावास या मर्यादित दिया जा सकता है, यह व्यायालय के विवेक पर निर्भर करेगा।

Flow: Non-English Text (Hindi) → Translate → Statute Prediction

The screenshot shows the LSI system's user interface after translating the Hindi input into English. The top navigation bar includes a "Deploy" button. The main content area is divided into several sections:

- Identify Statutes**: Shows "Step 2: Translating to English..." and "Step 3: Identifying Statutes...".
- Results**: Shows the "Input Language: Hindi | Input Type: Text Input".
- Intermediate Output**: Displays the English translation of the Hindi text, detailing a dispute between Amit Singh and Ramesh over land, leading to Ramesh's death.
- Statute Prediction Output**: Shows "Found 1 relevant statute sections (Score > 0.50):"
- | | Status Section | Prediction Score | Description |
|---|--------------------------------------|------------------|---|
| 0 | Section 302 in The Indian Penal Code | 0.6057 | Murder (Punishment) - Whoever commits murder shall be punished with death, or imprisonment for life, or |

Figure 9. Workflow for non-English input in the Legal Statute Identification (LSI) system. (Top) A user inputs a case summary in Hindi. (Bottom) The system first performs an intermediate translation of the Hindi text into English, then analyzes the translation to identify and output the most relevant statute, Section 302 (Murder) of the Indian Penal Code.

English Audio

The screenshot shows the 'Legal Statute Identification (LSI)' application. On the left, a sidebar titled 'Model Parameters' includes a 'Prediction Score Threshold' slider set to 0.50. The main area has a title 'Legal Statute Identification (LSI)' and a subtitle 'Predict relevant sections of The Indian Penal Code from text or audio input.' Below this, a section titled '1. Choose Input Type and Language' shows 'Text Input' selected over 'Audio Input (File Upload)'. An 'English' language dropdown is also present. A large section titled '2. Provide Input Data' contains a file upload interface where 'english-audio-1.mp3' (0.7MB) is being uploaded. A progress bar indicates 0.00 / 0.48. At the bottom right is a red button labeled 'Identify Statutes'.

The screenshot shows the results panel of the LSI system. It displays two steps: 'Step 1: Transcribing Audio...' and 'Step 3: Identifying Statutes...'. Below this is a 'Results' section with 'Input Language: English | Input Type: Audio Input (File Upload)'. The 'Intermediate Output' section contains a transcription of a news article about a journalist named Priya Malhotra. The 'Statute Prediction Output' section shows a table with one result:

Statute Section	Prediction Score	Description
Section 500 in The Indian Penal Code	0.5477	Defamation - Whoever defames another shall be punished with simple imprisonment for a term which may

Figure 10. Demonstration of the audio input module for the Legal Statute Identification (LSI) system. (Top) A user uploads an English audio file. (Bottom) The results panel shows the two-step process: first, an 'Intermediate Output' displaying the system's transcription of the audio, and second, the 'Statute Prediction Output' (Section 500 IPC for Defamation) based on the transcribed text.

Non-English Audio:

The figure consists of three vertically stacked screenshots of the LSI system's user interface.

Top Screenshot: This shows the initial step of the workflow. It features a sidebar titled "Model Parameters" with a "Prediction Score Threshold" slider set at 0.50. The main area is titled "Legal Statute Identification (LSI)" with the subtitle "Predict relevant sections of The Indian Penal Code from text or audio input." Below this, a section titled "1. Choose Input Type and Language" includes a radio button for "Text Input" (unchecked) and one for "Audio Input (File Upload)" (checked). A dropdown menu for "Select Input Language" shows "Hindi". A file upload interface displays a file named "hindi-audio-1.mp3" (0.9MB). A progress bar below it shows "0.00 / 1.00". A flow diagram at the bottom indicates the process: "Flow: Non-English Audio (Hindi) → Transcript → Translate → Statute Prediction". A large red button at the bottom right is labeled "Identify Statutes".

Middle Screenshot: This shows the system's processing steps. The sidebar remains the same. The main area has a title "Step 1: Transcribing Audio..." followed by "Step 2: Translating to English..." and "Step 3: Identifying Statutes...". Below this, a section titled "Results" shows "Input Language: Hindi | Input Type: Audio Input (File Upload)". Under "Intermediate Output", there is a transcription of a Hindi story about Sandeep and Mohit. Below it, an "English Translation (for LSI)" provides the context in English. Under "Statute Prediction Output", a green bar states "Found 1 relevant statute sections (Score > 0.50):".

Bottom Screenshot: This shows the final results. The sidebar and main area title remain the same. The "Statute Prediction Output" section displays a table with one row. The table has columns for "Statute Section" (Section 302 in The Indian Penal Code), "Prediction Score" (0.6478), and "Description" (Murder (Punishment) - Whoever commits murder shall be punished with death, or imprisonment for life, or).

Figure 11. Workflow for non-English audio input in the Legal Statute Identification (LSI) system.

(Top) A user uploads an audio file in Hindi. (Bottom) The results panel demonstrates the complete pipeline: an 'Intermediate Output' of the original Hindi transcription, followed by an

'English Translation (for LSI)', and concluding with the 'Statute Prediction Output' (Section 302 IPC) based on the translation.

7. CONCLUSION

This project proposed and demonstrated a Multilingual Legal Statute Recommendation System designed to assist India's judiciary, legal professionals, and citizens in efficiently identifying statutory provisions relevant to case descriptions. The system addresses critical challenges such as manual statute mapping, language barriers, and the lack of accessibility in rural courts.

Key findings and contributions include:

- **Speech-to-Text Integration:** Implemented Automatic Speech Recognition (ASR) using Whisper/IndicTTS to transcribe multilingual audio case inputs into text.
- **Multilingual Support:** Incorporated IndicBERT, MuRIL, and mBERT to detect and translate non-English case descriptions into standardized English, ensuring inclusivity across Hindi, Marathi, Tamil, Kannada, and English.
- **Legal Domain NLP:** Utilized InLegalBERT, a transformer model fine-tuned on Indian legal texts, combined with the LeSICiN dataset and citation networks to accurately recommend relevant statutes.
- **Ranked Statute Output:** Generated ranked lists of applicable statutes with confidence scores, improving interpretability and aiding quick decision-making.

Overall, the project demonstrates how Natural Language Processing, multilingual AI, and legal knowledge graphs can be combined to transform legal research and statutory interpretation. The framework enhances accessibility, speeds up legal aid, and provides significant value to lower courts and underrepresented communities.

By bridging the gap between speech, language diversity, and statutory complexity, this work contributes toward an intelligent, inclusive, and scalable legal recommendation system that supports equitable access to justice in India.

8. REFERENCES

- [1] Paul, S., Goyal, P., & Ghosh, S. (2022). Paul, S., Goyal, P., & Ghosh, S. (2022). LeSICiN: A heterogeneous graph-based approach for automatic legal statute identification from Indian legal documents. In Proceedings of the AAAI Conference on Artificial Intelligence, 36(11), 12699–12707. <https://doi.org/10.1609/aaai.v36i11.21404>
- [2] Bhattacharya, P., Ghosh, K., Pal, A., & Ghosh, S. (2020). Bhattacharya, P., Ghosh, K., Pal, A., & Ghosh, S. (2020). Hier-SPCNet: A legal statute hierarchy-based heterogeneous network for computing legal case document similarity. arXiv preprint arXiv:2007.03225. <https://arxiv.org/abs/2007.03225>
- [3] Chalkidis, I., Fergadiotis, M., Malakasiotis, P., Aletras, N., & Androutsopoulos, I. (2020). Chalkidis, I., Fergadiotis, M., Malakasiotis, P., Aletras, N., & Androutsopoulos, I. (2020). LEGAL-BERT: The Muppets straight out of law school. In Findings of the Association for Computational Linguistics: EMNLP 2020 (pp. 2898–2904). Association for Computational Linguistics. <https://aclanthology.org/2020.findings-emnlp.261/>
- [4] Paul, S., Mandal, A., Goyal, P., & Ghosh, S. (2023). Paul, S., Mandal, A., Goyal, P., & Ghosh, S. (2023). Pre-trained language models for the legal domain: A case study on Indian law. In Proceedings of the 19th International Conference on Artificial Intelligence and Law (pp. 195–199). ACM. <https://doi.org/10.1145/3594536.3595133>
- [5] Dhanani, J., Mehta, R., & Rana, D. (2022). Effective and Scalable Legal Judgment Recommendation Using Pre-learned Word Embedding. Complex & Intelligent Systems, 8, 3199–3213. <https://doi.org/10.1007/s40747-022-00673-1>

- [6] Ariai, F., Mackenzie, J., & Demartini, G. (2025). Natural Language Processing for the Legal Domain: A Survey of Tasks, Datasets, Models and Challenges. ACM Computing Surveys. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>
- [7] Joshi, A., Paul, S., Sharma, A., Goyal, P., Ghosh, S., & Modi, A. (2023). IL-TUR: Benchmark for Indian Legal Text Understanding and Reasoning. <https://exploration-lab.github.io/IL-TUR/>
- [8] Pandian, S., & Joshi, S. (2022). Autosuggestion of Relevant Cases and Statutes. <https://aclanthology.org/2022.findings-emnlp.831>
- [9] Feng, Y., Li, C., Ge, J., Luo, B., & Ng, V. (2020). Recommending Statutes: A Portable Method Based on Neural Networks. ACM Transactions on Knowledge Discovery from Data, 15(2), Article 16. <https://doi.org/10.1145/3424671>
- [10] Li, C., Ge, J., Cheng, K., Luo, B., & Chang, V. (2022). Statute Recommendation: Re-ranking Statutes by Modeling Case-Statute Relation with Interpretable Hand-Crafted Features. Information Sciences, 607, 1023–1040. <https://doi.org/10.1016/j.ins.2022.06.042>
- [11] Deshmukh, A., & Kamble, S. (2025). Deshmukh, A., & Kamble, S. (2025). IndianBailJudgments-1200: A Multi-Attribute Legal NLP Dataset for Bail Order Understanding in India. arXiv. <https://arxiv.org/abs/2507.02506>
- [12] Kalamkar, P., et al. (2022). Kalamkar, P., Agarwal, A., Tiwari, A., Gupta, S., Karn, S., & Raghavan, V. (2022). Named Entity Recognition in Indian court judgments. In Proceedings of the Natural Legal Language Processing Workshop 2022 (pp. 184–193). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2022.nllp-1.15>
- [13] Joshi, A., Pathak, N., Singh, K., & Khapra, M. M. (2023). INDICourtSumm: A Dataset for Judgment Summarization of Indian Court Cases. In Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (pp. 5863-5876). Association for Computational Linguistics. <https://aclanthology.org/2023.emnlp-main.362/>
- [14] Malik, S., et al. (2021). Malik, V., Sanjay, R., Nigam, S. K., Ghosh, K., Guha, S. K., Bhattacharya, A., & Modi, A. (2021). ILDC for CJPE: Indian Legal Documents Corpus for

Court Judgment Prediction and Explanation. In Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers) (pp. 4046–4062). Association for Computational Linguistics. <https://aclanthology.org/2021.acl-long.313/>

[15] Bhattacharya, M., & Ghosh, B. (2022). Shukla, A., Bhattacharya, P., Poddar, S., Mukherjee, R., Ghosh, K., Goyal, P., & Ghosh, S. (2022). Legal Case Document Summarization: Extractive and Abstractive Methods and their Evaluation. Proceedings of the 2nd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing (AAACL-IJCNLP) (pp. 1048–1064). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2022.aacl-main.77>